

```
In [1]: #####Assignment of Advanced Regression for house price predictions Using Ridge  
and Lasso algorithm
```

```
In [2]: # Importing libraries  
import os  
import sys  
  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn import linear_model  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Ridge  
from sklearn.linear_model import Lasso  
from sklearn.model_selection import GridSearchCV  
from sklearn.metrics import r2_score
```

```
In [3]: #reading the dataset
housePrice=pd.read_csv('train.csv')
housePrice.info
```

Out[3]: <bound method DataFrame.info of LotArea Street Alley LotShape \				Id	MSSubClass	MSZoning	LotFrontage
0	1	60	RL	65.0	8450	Pave	NaN
1	2	20	RL	80.0	9600	Pave	NaN
2	3	60	RL	68.0	11250	Pave	NaN
3	4	70	RL	60.0	9550	Pave	NaN
4	5	60	RL	84.0	14260	Pave	NaN
5	6	50	RL	85.0	14115	Pave	NaN
6	7	20	RL	75.0	10084	Pave	NaN
7	8	60	RL	NaN	10382	Pave	NaN
8	9	50	RM	51.0	6120	Pave	NaN
9	10	190	RL	50.0	7420	Pave	NaN
10	11	20	RL	70.0	11200	Pave	NaN
11	12	60	RL	85.0	11924	Pave	NaN
12	13	20	RL	NaN	12968	Pave	NaN
13	14	20	RL	91.0	10652	Pave	NaN
14	15	20	RL	NaN	10920	Pave	NaN
15	16	45	RM	51.0	6120	Pave	NaN
16	17	20	RL	NaN	11241	Pave	NaN
17	18	90	RL	72.0	10791	Pave	NaN
18	19	20	RL	66.0	13695	Pave	NaN
19	20	20	RL	70.0	7560	Pave	NaN
20	21	60	RL	101.0	14215	Pave	NaN
21	22	45	RM	57.0	7449	Pave	Grvl
22	23	20	RL	75.0	9742	Pave	NaN
23	24	120	RM	44.0	4224	Pave	NaN
24	25	20	RL	NaN	8246	Pave	NaN
25	26	20	RL	110.0	14230	Pave	NaN
26	27	20	RL	60.0	7200	Pave	NaN
27	28	20	RL	98.0	11478	Pave	NaN
28	29	20	RL	47.0	16321	Pave	NaN
29	30	30	RM	60.0	6324	Pave	NaN
...
1430	1431	60	RL	60.0	21930	Pave	NaN
1431	1432	120	RL	NaN	4928	Pave	NaN
1432	1433	30	RL	60.0	10800	Pave	Grvl
1433	1434	60	RL	93.0	10261	Pave	NaN
1434	1435	20	RL	80.0	17400	Pave	NaN
1435	1436	20	RL	80.0	8400	Pave	NaN
1436	1437	20	RL	60.0	9000	Pave	NaN
1437	1438	20	RL	96.0	12444	Pave	NaN
1438	1439	20	RM	90.0	7407	Pave	NaN
1439	1440	60	RL	80.0	11584	Pave	NaN
1440	1441	70	RL	79.0	11526	Pave	NaN
1441	1442	120	RM	NaN	4426	Pave	NaN
1442	1443	60	FV	85.0	11003	Pave	NaN
1443	1444	30	RL	NaN	8854	Pave	NaN
1444	1445	20	RL	63.0	8500	Pave	NaN
1445	1446	85	RL	70.0	8400	Pave	NaN
1446	1447	20	RL	NaN	26142	Pave	NaN
1447	1448	60	RL	80.0	10000	Pave	NaN
1448	1449	50	RL	70.0	11767	Pave	NaN
1449	1450	180	RM	21.0	1533	Pave	NaN
1450	1451	90	RL	60.0	9000	Pave	NaN
1451	1452	20	RL	78.0	9262	Pave	NaN
1452	1453	180	RM	35.0	3675	Pave	NaN
1453	1454	20	RL	90.0	17217	Pave	NaN

1454	1455	20	FV	62.0	7500	Pave	Pave	Reg
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
5	Lvl	AllPub	...	0	NaN	MnPrv	Shed	700	
6	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
7	Lvl	AllPub	...	0	NaN	NaN	Shed	350	
8	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
9	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
10	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
11	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
12	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
13	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
14	Lvl	AllPub	...	0	NaN	GdWo	NaN	0	
15	Lvl	AllPub	...	0	NaN	GdPrv	NaN	0	
16	Lvl	AllPub	...	0	NaN	NaN	Shed	700	
17	Lvl	AllPub	...	0	NaN	NaN	Shed	500	
18	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
19	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
20	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
21	Bnk	AllPub	...	0	NaN	GdPrv	NaN	0	
22	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
23	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
24	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
25	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
26	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
27	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
28	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
29	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
...	
1430	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1431	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1432	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1433	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1434	Low	AllPub	...	0	NaN	NaN	NaN	0	
1435	Lvl	AllPub	...	0	NaN	GdPrv	NaN	0	
1436	Lvl	AllPub	...	0	NaN	GdWo	NaN	0	
1437	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1438	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
1439	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1440	Bnk	AllPub	...	0	NaN	NaN	NaN	0	
1441	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1442	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1443	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1444	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1445	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1446	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1447	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

1448	Lvl	AllPub	...	0	NaN	GdWo	NaN	0
1449	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1450	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1451	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1452	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1453	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1454	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1455	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1456	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0
1457	Lvl	AllPub	...	0	NaN	GdPrv	Shed	2500
1458	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1459	Lvl	AllPub	...	0	NaN	NaN	NaN	0

	MoSold	YrSold	SaleType	SaleCondition	SalePrice			
0	2	2008	WD	Normal	208500			
1	5	2007	WD	Normal	181500			
2	9	2008	WD	Normal	223500			
3	2	2006	WD	Abnorml	140000			
4	12	2008	WD	Normal	250000			
5	10	2009	WD	Normal	143000			
6	8	2007	WD	Normal	307000			
7	11	2009	WD	Normal	200000			
8	4	2008	WD	Abnorml	129900			
9	1	2008	WD	Normal	118000			
10	2	2008	WD	Normal	129500			
11	7	2006	New	Partial	345000			
12	9	2008	WD	Normal	144000			
13	8	2007	New	Partial	279500			
14	5	2008	WD	Normal	157000			
15	7	2007	WD	Normal	132000			
16	3	2010	WD	Normal	149000			
17	10	2006	WD	Normal	90000			
18	6	2008	WD	Normal	159000			
19	5	2009	COD	Abnorml	139000			
20	11	2006	New	Partial	325300			
21	6	2007	WD	Normal	139400			
22	9	2008	WD	Normal	230000			
23	6	2007	WD	Normal	129900			
24	5	2010	WD	Normal	154000			
25	7	2009	WD	Normal	256300			
26	5	2010	WD	Normal	134800			
27	5	2010	WD	Normal	306000			
28	12	2006	WD	Normal	207500			
29	5	2008	WD	Normal	68500			
...
1430	7	2006	WD	Normal	192140			
1431	10	2009	WD	Normal	143750			
1432	8	2007	WD	Normal	64500			
1433	5	2008	WD	Normal	186500			
1434	5	2006	WD	Normal	160000			
1435	7	2008	COD	Abnorml	174000			
1436	5	2007	WD	Normal	120500			
1437	11	2008	New	Partial	394617			
1438	4	2010	WD	Normal	149700			
1439	11	2007	WD	Normal	197000			
1440	9	2008	WD	Normal	191000			
1441	5	2008	WD	Normal	149300			

1442	4	2009	WD	Normal	310000
1443	5	2009	WD	Normal	121000
1444	11	2007	WD	Normal	179600
1445	5	2007	WD	Normal	129000
1446	4	2010	WD	Normal	157900
1447	12	2007	WD	Normal	240000
1448	5	2007	WD	Normal	112000
1449	8	2006	WD	Abnorml	92000
1450	9	2009	WD	Normal	136000
1451	5	2009	New	Partial	287090
1452	5	2006	WD	Normal	145000
1453	7	2006	WD	Abnorml	84500
1454	10	2009	WD	Normal	185000
1455	8	2007	WD	Normal	175000
1456	2	2010	WD	Normal	210000
1457	5	2010	WD	Normal	266500
1458	4	2010	WD	Normal	142125
1459	6	2008	WD	Normal	147500

[1460 rows x 81 columns]>

In [4]: housePrice.head()

Out[4]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Util
0	1	60	RL	65.0	8450	Pave	NaN	Reg		Lvl Al
1	2	20	RL	80.0	9600	Pave	NaN	Reg		Lvl Al
2	3	60	RL	68.0	11250	Pave	NaN	IR1		Lvl Al
3	4	70	RL	60.0	9550	Pave	NaN	IR1		Lvl Al
4	5	60	RL	84.0	14260	Pave	NaN	IR1		Lvl Al

5 rows x 81 columns

In [5]: housePrice.shape

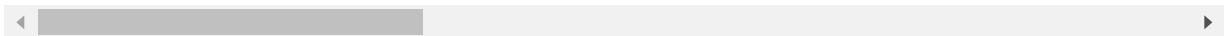
Out[5]: (1460, 81)

In [6]: housePrice.describe()

Out[6]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	Yea
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	5.575342	1971.200000
std	421.610009	42.300571	24.284752	9981.264932	1.382997	1.112799	30.200000
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.000000
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	1954.000000
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	1973.000000
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	2000.000000
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.000000

8 rows × 38 columns



In [7]: housePrice.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
Id                 1460 non-null int64
MSSubClass        1460 non-null int64
MSZoning          1460 non-null object
LotFrontage       1201 non-null float64
LotArea            1460 non-null int64
Street             1460 non-null object
Alley              91 non-null object
LotShape           1460 non-null object
LandContour       1460 non-null object
Utilities          1460 non-null object
LotConfig          1460 non-null object
LandSlope          1460 non-null object
Neighborhood       1460 non-null object
Condition1         1460 non-null object
Condition2         1460 non-null object
BldgType           1460 non-null object
HouseStyle         1460 non-null object
OverallQual        1460 non-null int64
OverallCond        1460 non-null int64
YearBuilt          1460 non-null int64
YearRemodAdd       1460 non-null int64
RoofStyle          1460 non-null object
RoofMatl           1460 non-null object
Exterior1st        1460 non-null object
Exterior2nd        1460 non-null object
MasVnrType         1452 non-null object
MasVnrArea         1452 non-null float64
ExterQual          1460 non-null object
ExterCond          1460 non-null object
Foundation         1460 non-null object
BsmtQual           1423 non-null object
BsmtCond           1423 non-null object
BsmtExposure       1422 non-null object
BsmtFinType1       1423 non-null object
BsmtFinSF1          1460 non-null int64
BsmtFinType2       1422 non-null object
BsmtFinSF2          1460 non-null int64
BsmtUnfSF          1460 non-null int64
TotalBsmtSF        1460 non-null int64
Heating             1460 non-null object
HeatingQC           1460 non-null object
CentralAir          1460 non-null object
Electrical          1459 non-null object
1stFlrSF            1460 non-null int64
2ndFlrSF            1460 non-null int64
LowQualFinSF        1460 non-null int64
GrLivArea           1460 non-null int64
BsmtFullBath        1460 non-null int64
BsmtHalfBath        1460 non-null int64
FullBath            1460 non-null int64
HalfBath             1460 non-null int64
BedroomAbvGr        1460 non-null int64
KitchenAbvGr        1460 non-null int64
KitchenQual          1460 non-null object
```

```

TotRmsAbvGrd      1460 non-null int64
Functional        1460 non-null object
Fireplaces         1460 non-null int64
FireplaceQu       770 non-null object
GarageType         1379 non-null object
GarageYrBlt       1379 non-null float64
GarageFinish       1379 non-null object
GarageCars          1460 non-null int64
GarageArea          1460 non-null int64
GarageQual         1379 non-null object
GarageCond         1379 non-null object
PavedDrive         1460 non-null object
WoodDeckSF         1460 non-null int64
OpenPorchSF        1460 non-null int64
EnclosedPorch      1460 non-null int64
3SsnPorch          1460 non-null int64
ScreenPorch         1460 non-null int64
PoolArea            1460 non-null int64
PoolQC              7 non-null object
Fence                281 non-null object
MiscFeature         54 non-null object
MiscVal              1460 non-null int64
MoSold              1460 non-null int64
YrSold              1460 non-null int64
SaleType             1460 non-null object
SaleCondition        1460 non-null object
SalePrice             1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

In [8]: *# Data Cleaning - Missing Value, Duplicates, Imputing, Dropping, Deleting, Exploration*

In [9]: *#Cleaning up variable Alley (Replacing NA => No Alley Access)*
`housePrice['Alley'].replace({np.nan:'No Alley Access'},inplace=True)
100*(housePrice['Alley'].value_counts()/housePrice['Alley'].count())`

Out[9]: No Alley Access 93.767123
Grvl 3.424658
Pave 2.808219
Name: Alley, dtype: float64

In [10]: *# As 94% of Alley is "No Alley access" it can be dropped*
`housePrice=housePrice.drop(['Alley'],axis=1)`

In [11]: `#Checking the dataset for the amount of nulls present
round(housePrice.isnull().sum()/len(housePrice.index),2).sort_values(ascending=False).head(18)`

Out[11]:

PoolQC	1.00
MiscFeature	0.96
Fence	0.81
FireplaceQu	0.47
LotFrontage	0.18
GarageType	0.06
GarageCond	0.06
GarageYrBlt	0.06
GarageFinish	0.06
GarageQual	0.06
BsmtFinType1	0.03
BsmtExposure	0.03
BsmtCond	0.03
BsmtQual	0.03
BsmtFinType2	0.03
MasVnrArea	0.01
MasVnrType	0.01
Exterior2nd	0.00
	dtype: float64

In [12]: `#Considering 10% as my threshold and dropping the column
round(housePrice.isnull().sum()/len(housePrice.index),2)[round(housePrice.isnull().sum()/len(housePrice.index),2).values>0.10]`

Out[12]:

LotFrontage	0.18
FireplaceQu	0.47
PoolQC	1.00
Fence	0.81
MiscFeature	0.96
	dtype: float64

In [13]: `housePrice = housePrice.drop(['LotFrontage', 'FireplaceQu', 'PoolQC', 'Fence', 'MiscFeature'], axis='columns')`

In [14]: `#verifying the columns for the missing values between 0-10%
round(housePrice.isnull().sum()/len(housePrice.index),2)[round(housePrice.isnull().sum()/len(housePrice.index),2).values>0.00]`

Out[14]:

MasVnrType	0.01
MasVnrArea	0.01
BsmtQual	0.03
BsmtCond	0.03
BsmtExposure	0.03
BsmtFinType1	0.03
BsmtFinType2	0.03
GarageType	0.06
GarageYrBlt	0.06
GarageFinish	0.06
GarageQual	0.06
GarageCond	0.06
	dtype: float64

In [15]: #convert the Year columns with the age to fill these columns with number

```
housePrice['YearBuiltOld'] = housePrice.YearBuilt.max()-housePrice.YearBuilt
housePrice['YearRemodAddOld'] = housePrice.YearRemodAdd.max()-housePrice.YearRemodAdd
housePrice['GarageYrBltOld'] = housePrice.GarageYrBlt.max()-housePrice.GarageYrBlt
housePrice['YrSoldOld'] = housePrice.YrSold.max()-housePrice.YrSold
housePrice[['YearBuilt', 'YearRemodAdd', 'GarageYrBlt', 'YrSold', 'YearBuiltOld', 'YearRemodAddOld', 'GarageYrBltOld', 'YrSoldOld']].sample(10)
```

Out[15]:

	YearBuilt	YearRemodAdd	GarageYrBlt	YrSold	YearBuiltOld	YearRemodAddOld	GarageYrBltOld
484	1962	2001	1963.0	2007	48	9	9
1015	2001	2001	2001.0	2009	9	9	9
1131	1991	1992	NaN	2007	19	18	18
406	1936	1950	1936.0	2008	74	60	60
1069	1949	2003	1985.0	2007	61	7	7
88	1915	1982	NaN	2009	95	28	28
174	1986	1986	1986.0	2008	24	24	24
592	1982	2003	1985.0	2008	28	7	7
205	1990	1990	1990.0	2009	20	20	20
651	1940	1950	1940.0	2009	70	60	60

In [16]: #Lets drop original Year columns

```
housePrice = housePrice.drop(['YearBuilt', 'YearRemodAdd', 'GarageYrBlt', 'YrSold'], axis='columns')
```

In [17]: #Imputing Missing values

```
housePrice.MasVnrType.fillna('None', inplace=True)
housePrice.MasVnrArea.fillna(housePrice.MasVnrArea.mean(), inplace=True)
housePrice.BsmtQual.fillna('TA', inplace=True)
housePrice.BsmtCond.fillna('TA', inplace=True)
housePrice.BsmtExposure.fillna('No', inplace=True)
housePrice.BsmtFinType1.fillna('Unf', inplace=True)
housePrice.BsmtFinType2.fillna('Unf', inplace=True)
housePrice.GarageType.fillna('Attchd', inplace=True)
housePrice.GarageFinish.fillna('Unf', inplace=True)
housePrice.GarageQual.fillna('TA', inplace=True)
housePrice.GarageCond.fillna('TA', inplace=True)
housePrice.GarageYrBltOld.fillna(-1, inplace=True)
```

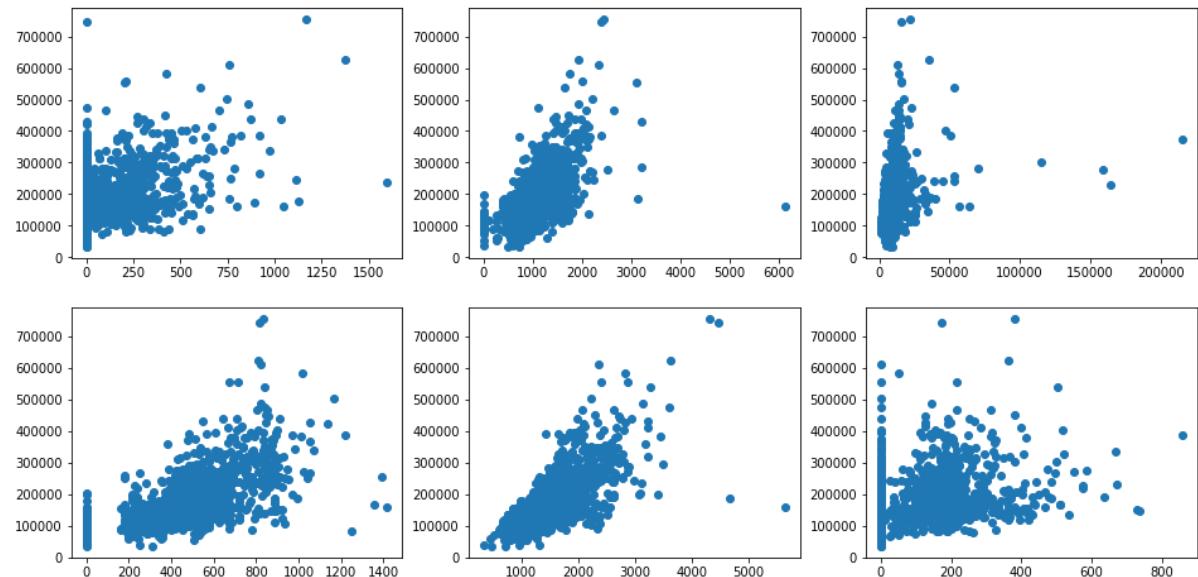
In [18]: #dropping Id, street and utilities as no impact

```
housePrice = housePrice.drop(['Id', 'Street', 'Utilities'], axis='columns')
```

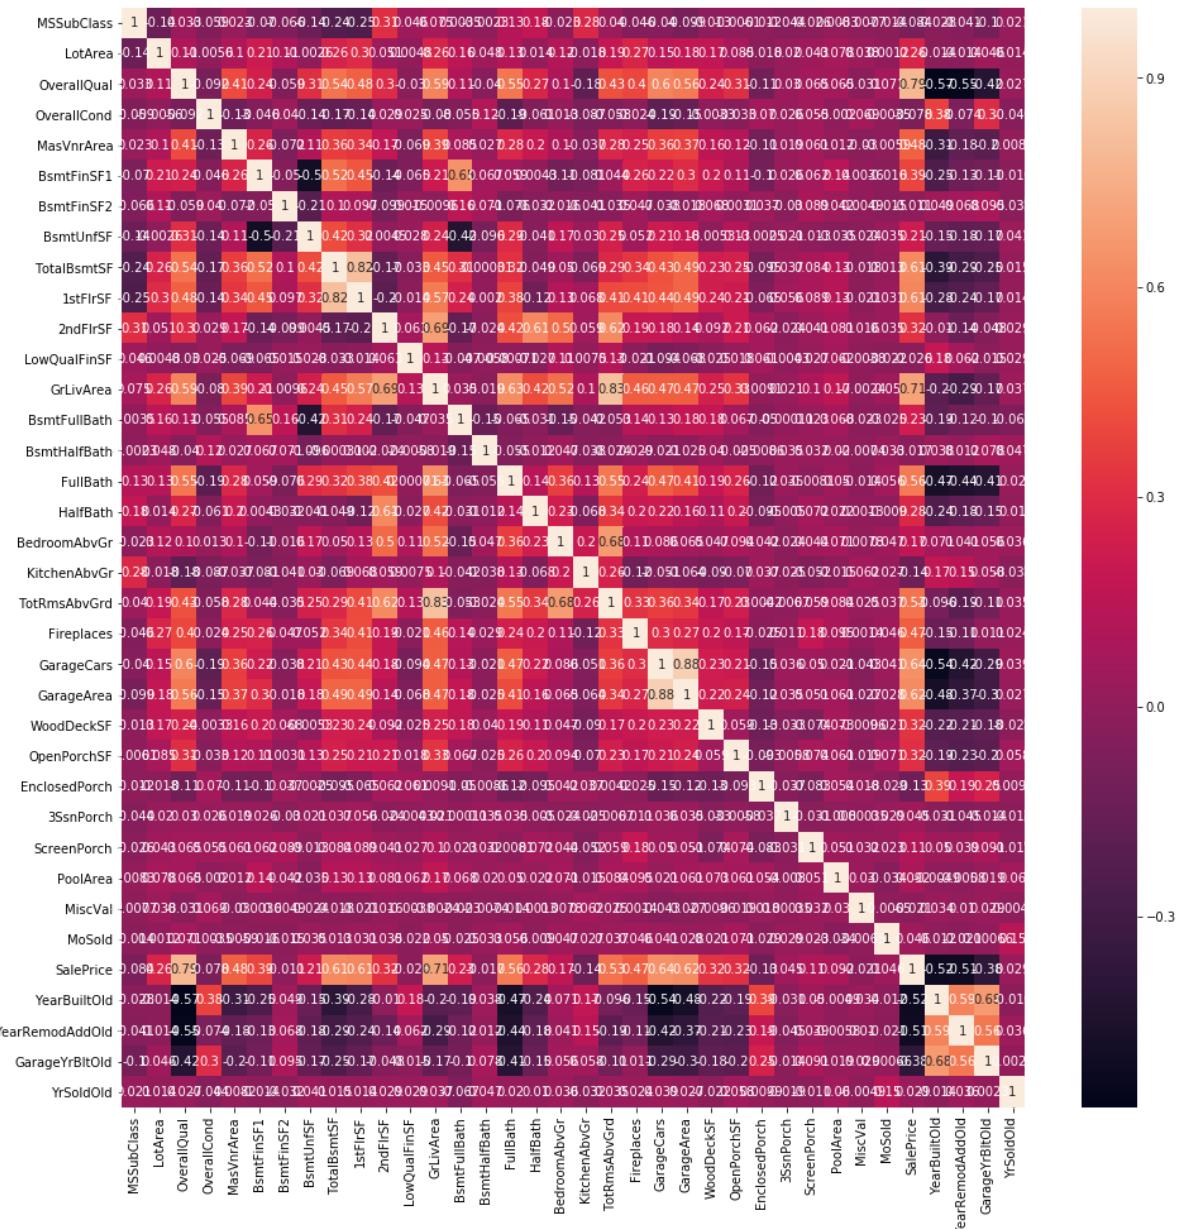
In [19]: #EDA Analysis with graphs

```
plt.figure(figsize=(16,8))
plt.subplot(2,3,1)
plt.scatter(housePrice.MasVnrArea,housePrice.SalePrice)
plt.subplot(2,3,2)
plt.scatter(housePrice.TotalBsmtSF,housePrice.SalePrice)
plt.subplot(2,3,3)
plt.scatter(housePrice[ 'LotArea '],housePrice.SalePrice)
plt.subplot(2,3,4)
plt.scatter(housePrice[ 'GarageArea '],housePrice.SalePrice)
plt.subplot(2,3,5)
plt.scatter(housePrice[ 'GrLivArea '],housePrice.SalePrice)
plt.subplot(2,3,6)
plt.scatter(housePrice[ 'WoodDeckSF '],housePrice.SalePrice)
```

Out[19]: <matplotlib.collections.PathCollection at 0x1d1e8b71d30>



```
In [20]: plt.figure(figsize=(16,16))
sns.heatmap(housePrice[list(housePrice.dtypes[housePrice.dtypes!= 'object'].index)].corr(), annot=True)
plt.show()
```



```
In [21]: #Handling the outliers by considering the Lower and upper quantile as 0.25 & 0.99  
housePrice.shape
```

Out[21]: (1460, 72)

```
In [22]: numCol = list(housePrice.dtypes[housePrice.dtypes != 'object'].index)
numCol = ['LotArea', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'TotalBsmtSF', '1stFlrSF', 'GrLivArea', 'OpenPorchSF',
          'EnclosedPorch', '3SsnPorch',
          'ScreenPorch', 'PoolArea', 'MiscVal', 'SalePrice']
def dropOutliers(x):
    list = []
    for col in numCol:
        Q1 = x[col].quantile(.25)
        Q3 = x[col].quantile(.75)
        IQR = Q3-Q1
        x = x[(x[col] >= (Q1-(1.5*IQR))) & (x[col] <= (Q3+(1.5*IQR)))]
    return x

housePrice = dropOutliers(housePrice)
```

```
In [23]: housePrice.shape
housePrice[list(housePrice.dtypes[housePrice.dtypes=='object'].index)].head()
```

Out[23]:

	MSZoning	LotShape	LandContour	LotConfig	LandSlope	Neighborhood	Condition1	Condition2	BldgType	PavedDrive	TotalBsmtSF	1stFlrSF	2ndFlrSF	Unf	Fin	ExterQual	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinType2	H
0	RL	Reg		Lvl	Inside	Gtl	CollgCr	Norm	Residential	Paved	0	850	0	N	Ex	TA	AsphSeal	TA	No	GLQ	Unf	Y
1	RL	Reg		Lvl	FR2	Gtl	Veenker	Feedr	Residential	Paved	0	800	0	N	Ex	TA	AsphSeal	FR2	Fr	ALQ	Unf	Y
2	RL	IR1		Lvl	Inside	Gtl	CollgCr	Norm	Residential	Paved	0	800	0	N	Ex	TA	AsphSeal	TA	Fr	GLQ	Unf	Y
3	RL	IR1		Lvl	Corner	Gtl	Crawfor	Norm	Residential	Paved	0	800	0	N	Ex	TA	AsphSeal	FR2	Fr	ALQ	Unf	Y
4	RL	IR1		Lvl	FR2	Gtl	NoRidge	Norm	Residential	Paved	0	800	0	N	Ex	TA	AsphSeal	FR2	Fr	GLQ	Unf	Y

5 rows × 36 columns

```
In [24]: #Below columns have some kind of order and hence check ordinal in nature
housePrice[['LandSlope', 'ExterQual', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'HeatingQC',
            'CentralAir', 'KitchenQual', 'GarageFinish', 'GarageQual', 'GarageCond', 'ExterCond', 'LotShape']].head()
```

Out[24]:

	LandSlope	ExterQual	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinType2	H
0	Gtl	Gd	Gd	TA	No	GLQ	Unf	Y
1	Gtl	TA	Gd	TA	Gd	ALQ	Unf	Y
2	Gtl	Gd	Gd	TA	Mn	GLQ	Unf	Y
3	Gtl	TA	TA	Gd	No	ALQ	Unf	Y
4	Gtl	Gd	Gd	TA	Av	GLQ	Unf	Y

```
In [25]: housePrice['LandSlope'] = housePrice.LandSlope.map({'Gtl':0,'Mod':1,'Sev':2})
housePrice['ExterQual'] = housePrice.ExterQual.map({'Po':0,'Fa':1,'TA':2,'Gd':3,'Ex':4})
housePrice['BsmtQual'] = housePrice.BsmtQual.map({'NA':0,'Po':1,'Fa':2,'TA':3,'Gd':4,'Ex':5})
housePrice['BsmtCond'] = housePrice.BsmtCond.map({'NA':0,'Po':1,'Fa':2,'TA':3,'Gd':4,'Ex':5})
housePrice['BsmtExposure'] = housePrice.BsmtExposure.map({'NA':0,'No':1,'Mn':2,'Av':3,'Gd':4})
housePrice['BsmtFinType1'] = housePrice.BsmtFinType1.map({'NA':0,'Unf':1,'LwQ':2,'Rec':3,'BLQ':4,'ALQ':5,'GLQ':6})
housePrice['BsmtFinType2'] = housePrice.BsmtFinType2.map({'NA':0,'Unf':1,'LwQ':2,'Rec':3,'BLQ':4,'ALQ':5,'GLQ':6})
housePrice['HeatingQC'] = housePrice.HeatingQC.map({'Po':0,'Fa':1,'TA':2,'Gd':3,'Ex':4})
housePrice['CentralAir'] = housePrice.CentralAir.map({'N':0,'Y':1})
housePrice['KitchenQual'] = housePrice.KitchenQual.map({'Po':0,'Fa':1,'TA':2,'Gd':3,'Ex':4})
housePrice['GarageFinish'] = housePrice.GarageFinish.map({'NA':0,'Unf':1,'RFn':2,'Fin':3})
housePrice['GarageQual'] = housePrice.GarageQual.map({'NA':0,'Po':1,'Fa':2,'TA':3,'Gd':4,'Ex':5})
housePrice['GarageCond'] = housePrice.GarageCond.map({'NA':0,'Po':1,'Fa':2,'TA':3,'Gd':4,'Ex':5})
housePrice['ExterCond'] = housePrice.ExterCond.map({'Po':0,'Fa':1,'TA':2,'Gd':3,'Ex':4})
housePrice['LotShape'] = housePrice.LotShape.map({'IR1':0,'IR2':1,'IR3':2,'Reg':3})
```

```
In [26]: #check converted columns
housePrice[['LandSlope','ExterQual','BsmtQual','BsmtCond','BsmtExposure','BsmtFinType1','BsmtFinType2','HeatingQC','CentralAir','KitchenQual','GarageFinish','GarageQual','GarageCond','ExterCond','LotShape']].head()
```

Out[26]:

	LandSlope	ExterQual	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinType2	H
0	0	3	4	3	1	6		1
1	0	2	4	3	4	5		1
2	0	3	4	3	2	6		1
3	0	2	3	4	1	5		1
4	0	3	4	3	3	6		1

In [27]: #Filling dummy column with the actual dataset

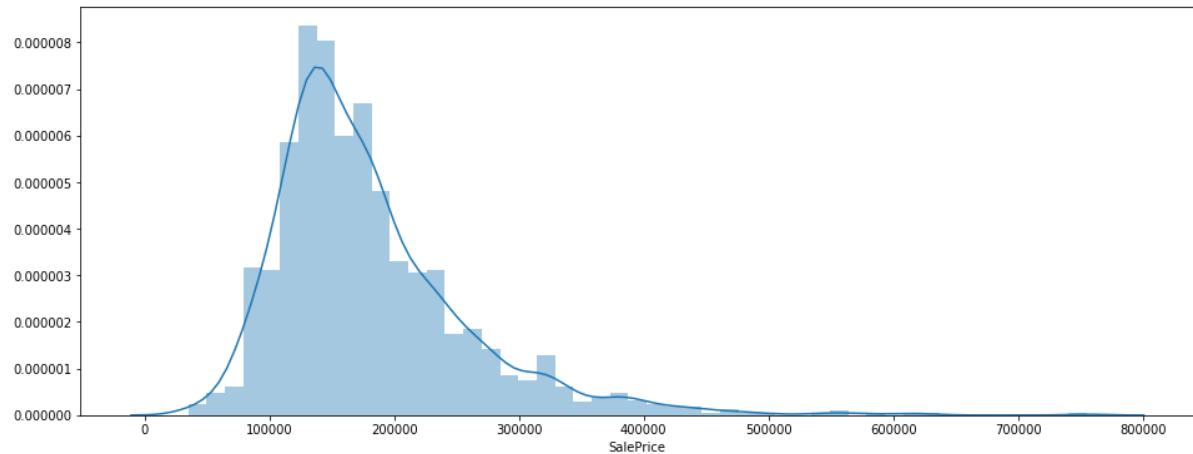
```
dummyCol = pd.get_dummies(housePrice[['MSZoning', 'LandContour', 'LotConfig', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'Foundation', 'Heating', 'Electrical', 'Functional', 'GarageType', 'PavedDrive', 'SaleType', 'SaleCondition']], drop_first=True)

housePrice = pd.concat([housePrice, dummyCol], axis='columns')

housePrice = housePrice.drop(['MSZoning', 'LandContour', 'LotConfig', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'Foundation', 'Heating', 'Electrical', 'Functional', 'GarageType', 'PavedDrive', 'SaleType', 'SaleCondition'], axis='columns')
```

In [28]: #verify the distribution of target variable

```
plt.figure(figsize=(16,6))
sns.distplot(housePrice.SalePrice)
plt.show()
```



```
In [29]: #Creating train and test dataset for verification
from sklearn.model_selection import train_test_split
df_train,df_test = train_test_split(housePrice,train_size=0.7,test_size=0.3,random_state=42)

housePrice[['LandSlope','ExterQual','BsmtQual','BsmtCond','BsmtExposure','BsmtFinType1','BsmtFinType2','HeatingQC',
           'CentralAir', 'KitchenQual','GarageFinish','GarageQual','GarageCond','ExterCond','LotShape']].head()
```

Out[29]:

	LandSlope	ExterQual	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	BsmtFinType2	H
0	0	3	4	3	1	6		1
1	0	2	4	3	4	5		1
2	0	3	4	3	2	6		1
3	0	2	3	4	1	5		1
4	0	3	4	3	3	6		1



In [30]: #Scaling the train dataset with dependent variable

```
from sklearn.preprocessing import StandardScaler

numCol = ['MSSubClass', 'LotArea', 'OverallQual', 'OverallCond', 'MasVnrArea', 'Bsm
tFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
          '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'Bsm
tHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr',
          'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageCars', 'GarageArea'
, 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch',
          '3SsnPorch',
          'ScreenPorch', 'PoolArea', 'MiscVal', 'SalePrice']

scaler = StandardScaler()
df_train[numCol] = scaler.fit_transform(df_train[numCol])
df_test[numCol] = scaler.transform(df_test[numCol])
```

E:\Python\lib\site-packages\ipykernel_launcher.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
# This is added back by InteractiveShellApp.init_path()
```

E:\Python\lib\site-packages\pandas\core\indexing.py:543: SettingWithCopyWarni
ng:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
self.obj[item] = s
```

E:\Python\lib\site-packages\ipykernel_launcher.py:12: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
if sys.path[0] == '':
```

E:\Python\lib\site-packages\pandas\core\indexing.py:543: SettingWithCopyWarni
ng:

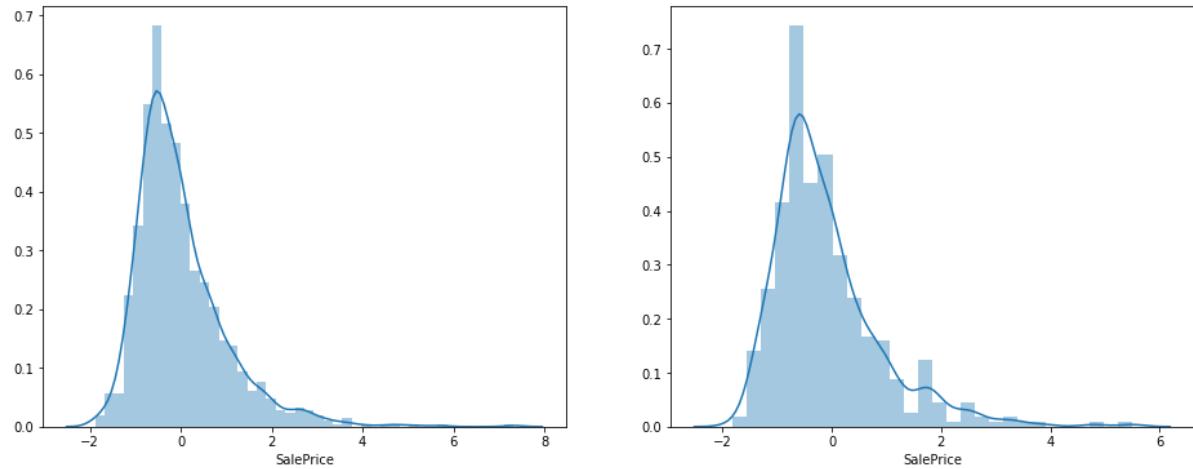
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
self.obj[item] = s
```

```
In [31]: #check the distribution
plt.figure(figsize=(16,6))
plt.subplot(121)
sns.distplot(df_train.SalePrice)
plt.subplot(122)
sns.distplot(df_test.SalePrice)
```

Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1d1e9f91e48>



```
In [32]: #Spliting the dependent and independent variable
y_train = df_train.pop('SalePrice')
X_train = df_train
```

```
y_test = df_test.pop('SalePrice')
X_test = df_test
```

In [36]: #Model Building with Ridge

```
params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1,
                   0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                   4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500, 1000 ]}

ridge = Ridge()

# cross validation
folds = 5
model_cv = GridSearchCV(estimator = ridge,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)
model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed:    1.7s finished
```

Out[36]: GridSearchCV(cv=5, error_score='raise-deprecating',
estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
max_iter=None, normalize=False, random_state=None,
solver='auto', tol=0.001),
iid='warn', n_jobs=None,
param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
100, 500, 1000]}},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring='neg_mean_absolute_error', verbose=1)

```
In [47]: cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results = cv_results[cv_results['param_alpha'] <= 1000]
cv_results.head()
```

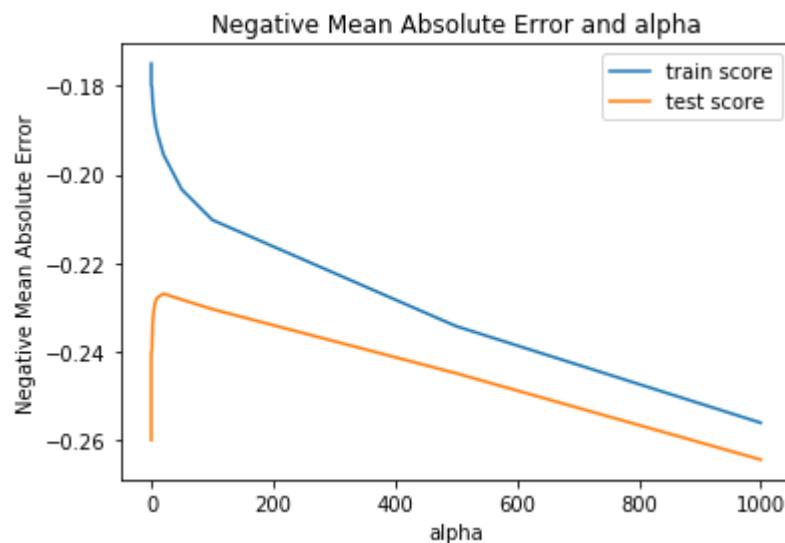
Out[47]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_te
0	0.011098	0.002158	0.002798	0.000749	0.0001	{'alpha': 0.0001}	-
1	0.009395	0.001958	0.002002	0.000898	0.001	{'alpha': 0.001}	-
2	0.009195	0.001165	0.001799	0.000400	0.01	{'alpha': 0.01}	-
3	0.009994	0.000632	0.002198	0.000400	0.05	{'alpha': 0.05}	-
4	0.008995	0.001095	0.001798	0.000400	0.1	{'alpha': 0.1}	-

5 rows × 21 columns

```
# plotting mean test and train scoes with alpha
cv_results['param_alpha'] = cv_results['param_alpha'].astype('int32')

# plotting
plt.plot(cv_results['param_alpha'], cv_results['mean_train_score'])
plt.plot(cv_results['param_alpha'], cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.ylabel('Negative Mean Absolute Error')
plt.title("Negative Mean Absolute Error and alpha")
plt.legend(['train score', 'test score'], loc='upper right')
plt.show()
```



```
In [49]: model_cv.best_params_
```

Out[49]: {'alpha': 20}

```
In [50]: alpha = 20
ridge = Ridge(alpha=alpha)

ridge.fit(X_train, y_train)
#Predictor Variables from the Model built using Ridge Regression:
ridge.coef_
```

```
Out[50]: array([-6.94874307e-02,  6.83512111e-02,  1.91163796e-02, -1.02702444e-02,
   1.69177621e-01,  7.59735368e-02,  9.37878124e-02,  1.00951315e-01,
  -1.11529249e-02,  8.27290150e-02, -1.35542545e-02,  6.58750835e-02,
  1.60753428e-03,  9.21320034e-02,  1.17982760e-02,  5.08248498e-03,
 -1.53962364e-02,  8.21954396e-02,  4.63484144e-03, -3.54677093e-02,
  6.39963352e-02,  1.46891973e-01,  1.08964374e-02,  1.72539399e-01,
  2.06440938e-02, -5.02669431e-03,  2.59694899e-02,  1.97819805e-02,
 -5.03812058e-02, -4.04939921e-02,  8.95019648e-02,  7.21267497e-02,
  2.40734374e-02,  1.36170197e-02,  4.83059580e-03,  6.85625925e-02,
  6.16986069e-02, -2.45860787e-02,  2.28999932e-02, -1.30097967e-02,
  1.35028841e-02,  7.58619451e-03,  3.36025838e-02,  0.00000000e+00,
 -4.07271359e-03, -1.02827315e-02, -1.85121795e-03, -1.64410803e-04,
 -9.60771185e-04,  6.91584564e-03,  2.08214420e-02,  3.49275831e-02,
  4.82259423e-02, -3.18697276e-02,  8.06871322e-02, -3.47126628e-02,
  5.05855973e-02,  1.10930858e-01, -4.70889857e-02, -7.75907203e-03,
 -1.57394740e-02,  2.45608543e-04,  3.43221032e-02,  7.19757048e-02,
 -3.45941629e-02, -1.09112253e-01,  1.48506586e-01, -1.41909483e-01,
 -8.02841881e-02, -2.90214414e-02,  1.74144431e-02, -8.26559908e-02,
 -6.93283902e-02,  2.67996594e-02, -7.00582810e-02,  1.61384230e-01,
  2.54346495e-01, -2.98807221e-02, -2.73171305e-02, -6.04124333e-03,
 -6.93813240e-02,  1.30055933e-02,  1.48670146e-01, -7.71950019e-02,
 -4.07863977e-03,  1.82166076e-02,  1.03919457e-01, -1.67682756e-02,
 -1.27841279e-01, -5.04813786e-02,  5.93816353e-02, -1.16641606e-02,
  1.33847424e-02,  1.63602130e-02,  1.41462954e-01,  0.00000000e+00,
 -1.93899300e-01,  9.87957071e-03,  1.33846667e-02,  3.88079377e-02,
  2.21131648e-02, -3.56898732e-02, -4.93955317e-02,  3.34636235e-02,
  7.34048485e-02, -3.22792200e-02, -4.65354447e-02, -3.04358903e-02,
  4.93346628e-03,  2.13837511e-02,  4.29644975e-03, -1.23927570e-03,
  3.51990874e-02,  1.21160654e-03,  3.02273288e-03,  1.29004201e-02,
  1.10227925e-02, -1.33936957e-02, -7.47487025e-02, -3.46888385e-02,
  2.06887943e-01,  1.87261643e-03,  2.33314919e-03,  1.28730853e-01,
 -3.04277590e-03, -3.33460311e-02, -1.66277311e-02,  0.00000000e+00,
 -3.90708220e-03, -2.15745888e-02, -1.50192590e-02,  2.49182424e-02,
  1.59202310e-02, -7.76094444e-02,  9.57879261e-03, -6.68078338e-03,
  2.06605100e-02,  2.27103059e-02, -3.04277590e-03, -3.39037942e-02,
 -2.71061895e-02,  3.82163280e-03, -1.79230350e-03,  1.25930547e-04,
 -5.63756551e-03,  1.62524418e-02,  1.16859136e-02, -6.85104120e-03,
  5.16772016e-02, -2.72520870e-02, -1.34547869e-02,  8.02323895e-02,
  2.46876453e-02, -1.76462878e-02,  3.55264003e-02,  7.65662851e-02,
 -1.93195656e-02, -1.05262778e-02,  4.67688370e-04, -1.47557056e-02,
  8.71171730e-04,  0.00000000e+00,  1.34168455e-02,  2.02123572e-02,
  1.38968808e-02,  3.57640879e-03, -2.73266363e-02, -9.64010069e-03,
 -2.73774393e-02, -4.08590395e-03, -1.51603097e-02, -2.97011907e-02,
  1.16491645e-01,  1.48724747e-02, -5.25909620e-03,  5.87908978e-03,
 -7.44720566e-04,  2.80639578e-02, -1.15113633e-03, -4.93349148e-03,
 -7.51490023e-03,  0.00000000e+00,  3.30093172e-02, -3.24694050e-02,
 -1.06579328e-02,  1.17115657e-01,  2.12358070e-02, -1.76128421e-03,
  1.73514542e-02,  2.28541527e-02, -1.95437570e-02,  5.10665855e-02,
  6.91423986e-02])
```

```
In [53]: #R-squared value of test and train data
from sklearn import metrics
y_train_pred = ridge.predict(X_train)
print(metrics.r2_score(y_true=y_train, y_pred=y_train_pred))
```

0.9031439595017332

```
In [54]: #####Best alpha value for Ridge : {'alpha': 0.9}
#Build model with Lasso
lasso = Lasso()

# cross validation
model_cv = GridSearchCV(estimator = lasso,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)

model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

E:\Python\lib\site-packages\sklearn\linear_model\coordinate_descent.py:475: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.13920341294216598, tolerance: 0.0834479486660015

positive)

[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed: 2.6s finished

```
Out[54]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                      max_iter=1000, normalize=False, positive=False,
                                      precompute=False, random_state=None,
                                      selection='cyclic', tol=0.0001, warm_start=False),
                      iid='warn', n_jobs=None,
                      param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
                                           0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                                           4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
                                           100, 500, 1000]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_absolute_error', verbose=1)
```

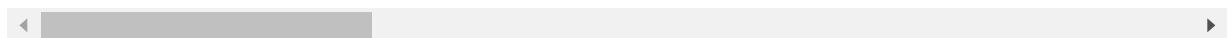
```
In [55]: cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results
```

Out[55]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_t
0	0.147509	3.312605e-02	0.001799	3.998281e-04	0.0001	{'alpha': 0.0001}	
1	0.075354	1.395784e-02	0.001599	4.898045e-04	0.001	{'alpha': 0.001}	
2	0.015390	1.958292e-03	0.001799	7.478695e-04	0.01	{'alpha': 0.01}	
3	0.012792	3.865033e-03	0.001999	6.322590e-04	0.05	{'alpha': 0.05}	
4	0.008795	1.165215e-03	0.001399	4.897269e-04	0.1	{'alpha': 0.1}	
5	0.008395	4.888698e-04	0.001599	4.892985e-04	0.2	{'alpha': 0.2}	
6	0.008395	1.019244e-03	0.001599	8.003594e-04	0.3	{'alpha': 0.3}	
7	0.007196	9.786937e-04	0.001599	4.890259e-04	0.4	{'alpha': 0.4}	
8	0.006596	4.897466e-04	0.001599	4.895324e-04	0.5	{'alpha': 0.5}	
9	0.007795	7.482523e-04	0.001998	6.321838e-04	0.6	{'alpha': 0.6}	
10	0.007195	7.479080e-04	0.001998	6.316563e-04	0.7	{'alpha': 0.7}	
11	0.007396	7.995849e-04	0.001798	7.484049e-04	0.8	{'alpha': 0.8}	
12	0.007595	7.997277e-04	0.001799	3.998284e-04	0.9	{'alpha': 0.9}	
13	0.007195	7.475518e-04	0.001798	3.995899e-04	1	{'alpha': 1.0}	
14	0.006996	6.319574e-04	0.001599	4.896873e-04	2	{'alpha': 2.0}	
15	0.007196	7.477679e-04	0.001598	4.894733e-04	3	{'alpha': 3.0}	
16	0.007397	4.887729e-04	0.001598	4.893175e-04	4	{'alpha': 4.0}	
17	0.007396	7.985843e-04	0.001798	3.995449e-04	5	{'alpha': 5.0}	
18	0.006797	7.481499e-04	0.001798	3.995672e-04	6	{'alpha': 6.0}	
19	0.006196	3.990416e-04	0.001798	3.996617e-04	7	{'alpha': 7.0}	
20	0.006796	7.479332e-04	0.001398	4.894731e-04	8	{'alpha': 8.0}	
21	0.006796	3.996373e-04	0.000999	1.168008e-07	9	{'alpha': 9.0}	
22	0.005997	5.309834e-07	0.001198	3.995420e-04	10	{'alpha': 10.0}	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_t
23	0.006196	3.991366e-04	0.001799	3.998779e-04	20	{'alpha': 20}	
24	0.005997	2.780415e-07	0.001198	3.992560e-04	50	{'alpha': 50}	
25	0.006796	3.998050e-04	0.001598	4.891238e-04	100	{'alpha': 100}	
26	0.006597	7.995964e-04	0.001198	3.993280e-04	500	{'alpha': 500}	
27	0.006996	8.941504e-04	0.001399	4.886566e-04	1000	{'alpha': 1000}	

28 rows × 21 columns



In [56]: #R-squared value

```
model_cv1 = GridSearchCV(estimator = lasso,
                         param_grid = params,
                         scoring= 'r2',
                         cv = folds,
                         verbose = 1,
                         return_train_score=True)

# fit the model
model_cv1.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

E:\Python\lib\site-packages\sklearn\linear_model\coordinate_descent.py:475: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 0.13920341294216598, tolerance: 0.0834479486660015
positive)

[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed: 2.8s finished

Out[56]: GridSearchCV(cv=5, error_score='raise-deprecating', estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,

```
max_iter=1000, normalize=False, positive=False, precompute=False, random_state=None, selection='cyclic', tol=0.0001, warm_start=False),
iid='warn', n_jobs=None, param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500, 1000]}, pre_dispatch='2*n_jobs', refit=True, return_train_score=True, scoring='r2', verbose=1)
```

```
In [57]: cv_results1 = pd.DataFrame(model_cv1.cv_results_)
```

Out[57]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_t
0	0.144710	6.498708e-02	0.002399	4.899015e-04	0.0001	{'alpha': 0.0001}	
1	0.071556	1.157906e-02	0.001998	3.504023e-07	0.001	{'alpha': 0.001}	
2	0.013792	1.600295e-03	0.001799	3.997327e-04	0.01	{'alpha': 0.01}	
3	0.010394	1.742979e-03	0.001999	6.324853e-04	0.05	{'alpha': 0.05}	
4	0.011794	1.325800e-03	0.002997	2.463915e-06	0.1	{'alpha': 0.1}	
5	0.011193	7.475394e-04	0.002398	4.894150e-04	0.2	{'alpha': 0.2}	
6	0.011193	3.997335e-04	0.002399	4.895706e-04	0.3	{'alpha': 0.3}	
7	0.008195	1.165280e-03	0.001799	7.479460e-04	0.4	{'alpha': 0.4}	
8	0.008795	1.164560e-03	0.001798	3.994466e-04	0.5	{'alpha': 0.5}	
9	0.009394	1.018552e-03	0.002398	7.992510e-04	0.6	{'alpha': 0.6}	
10	0.008595	4.895901e-04	0.002199	7.479210e-04	0.7	{'alpha': 0.7}	
11	0.010394	7.993107e-04	0.002199	4.002097e-04	0.8	{'alpha': 0.8}	
12	0.010194	9.793940e-04	0.002198	3.998520e-04	0.9	{'alpha': 0.9}	
13	0.006996	6.317313e-04	0.001798	3.994942e-04	1	{'alpha': 1.0}	
14	0.007395	1.019356e-03	0.001599	4.894537e-04	2	{'alpha': 2.0}	
15	0.006596	4.883268e-04	0.001798	3.996376e-04	3	{'alpha': 3.0}	
16	0.007795	9.781292e-04	0.001640	5.290908e-04	4	{'alpha': 4.0}	
17	0.006996	6.318821e-04	0.001598	4.891618e-04	5	{'alpha': 5.0}	
18	0.007595	1.019048e-03	0.001599	4.898237e-04	6	{'alpha': 6.0}	
19	0.006397	7.992745e-04	0.001598	4.896874e-04	7	{'alpha': 7.0}	
20	0.006797	7.470669e-04	0.001598	4.885974e-04	8	{'alpha': 8.0}	
21	0.006196	3.995181e-04	0.001399	4.895705e-04	9	{'alpha': 9.0}	
22	0.006196	3.994704e-04	0.001599	4.894926e-04	10	{'alpha': 10.0}	

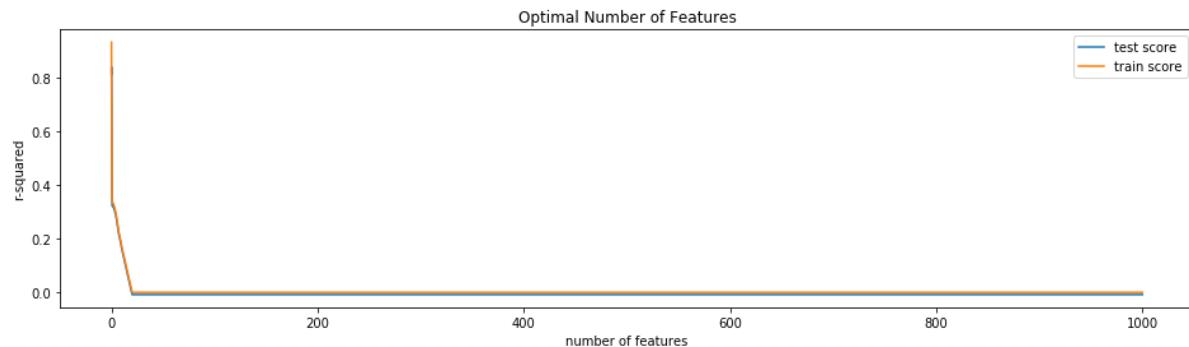
	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_t
23	0.006597	7.996321e-04	0.001398	4.893370e-04	20	{'alpha': 20}	
24	0.006596	4.902713e-04	0.001798	3.996856e-04	50	{'alpha': 50}	
25	0.005997	3.371748e-07	0.001598	4.898236e-04	100	{'alpha': 100}	
26	0.007196	7.477295e-04	0.001598	4.901165e-04	500	{'alpha': 500}	
27	0.006196	3.986837e-04	0.001399	4.894343e-04	1000	{'alpha': 1000}	

28 rows × 21 columns

```
In [58]: # plotting cv results
plt.figure(figsize=(16,4))
```

```
plt.plot(cv_results1["param_alpha"], cv_results1["mean_test_score"])
plt.plot(cv_results1["param_alpha"], cv_results1["mean_train_score"])
plt.xlabel('number of features')
plt.ylabel('r-squared')
plt.title("Optimal Number of Features")
plt.legend(['test score', 'train score'], loc='upper right')
```

```
Out[58]: <matplotlib.legend.Legend at 0x1d1e9422f60>
```



```
In [60]: #value of optimum number of parameters
print(model_cv.best_params_)
print(model_cv.best_score_)
```

```
{'alpha': 0.001}
-0.22686016003961607
```

```
In [ ]: ###Best alpha value for Lasso : {'alpha': 0.001}
###Best alpha value for Ridge : {'alpha': 0.9}
```